

System For Pratical Evaluations of Network Administration Course

Diogo Nunes

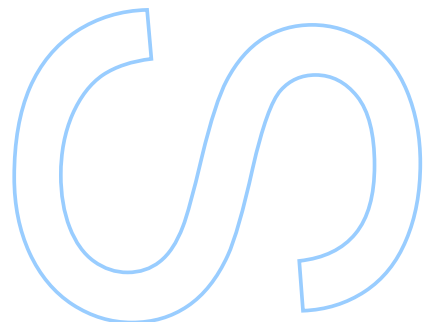
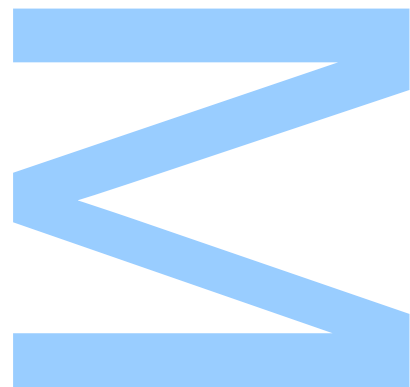
Mestrado em Engenharia de Redes e Sistemas Informáticos

Departamento de Ciência de Computadores

2025

Orientador

Prof. Rui Prior, Faculdade de Ciências

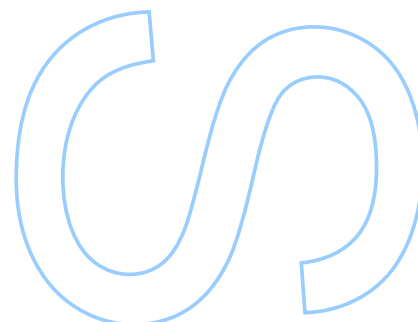




Todas as correções determinadas
pelo júri, e só essas, foram efetuadas.

O Presidente do Júri,

Porto, ____/____/____



Acknowledgements

Acknowledge ALL the people! ☐

Resumo

Este tese é sobre alguma coisa

Palavras-chave: física (keywords em português)

Abstract

This thesis is about something, I guess.

Keywords: Computer Science

Table of Contents

List of Figures.....	vii
List of Abbreviations.....	viii
1. Introduction.....	1
1.1. Problem Statement.....	1
1.1.1. Need for Automated Evaluation.....	1
1.1.2. Limitations of Current Solutions.....	2
1.2. Aims and Objectives.....	2
2. Background.....	4
2.1. Overview of Used Technologies.....	4
2.1.1. Virtualization.....	4
2.1.2. GNS3.....	5
2.1.3. Architecture.....	5
2.1.3.1. Controller.....	5
2.1.3.2. Compute.....	6
2.1.3.3. GUI.....	6
2.1.4. Proxmox VE.....	6
2.1.4.1. Virtualization Technologies.....	7
2.1.4.2. KVM.....	8
2.1.4.3. LXC.....	8
2.1.5. LDAP.....	8
2.2. Virtualized Lab Environments.....	9
2.3. Python Web Frameworks for API-Based Systems.....	9
2.3.1. Python.....	9
2.3.2. WSGI.....	9
2.3.2.1. Flask.....	11
2.3.3. ASGI.....	11
2.3.3.1. FastAPI.....	12
2.4. Long running task processing approaches.....	13
2.4.1. The need for asynchronous processing in API-heavy applications.....	13
2.4.2. Asyncio.....	13
2.4.3. Celery.....	14
2.5. System administration automation tools.....	14
2.5.1. Nornir.....	14
2.5.2. Ansible.....	15

3. Related Work.....	16
3.1. Programming Evaluation Systems	16
3.1.1. Mooshak and lessons learned	17
3.2. Cisco Packet Tracer.....	18
4. System Architecture & Design	19
4.1. Functional Use Cases.....	19
4.1.1. System Supports Both Directory Integration and Local Authentication ..	19
4.1.2. Teacher Creates a New Exercise	20
4.1.3. System Supports Enrollment of Users into Exercises	20
4.1.4. Student Starts Their Exercise VM.....	21
4.1.5. Student Submits Solution.....	21
4.1.6. Teacher Reviews Student Results	21
4.1.7. Student Views Exercise Feedback	22
4.2. System Architecture Overview	22
4.3. Proxmox VE	23
4.3.1. Why Proxmox VE?	23
4.3.2. Proxmox VE Limitations.....	24
4.3.3. Proxmox VE Firewall.....	24
4.3.4. Maximing resource usage.....	25
4.3.4.1. Exploration of containers as a full substitute for VMs	25
4.3.5. VM Lifecycle	26
4.4. GNS3	26
4.5. High-level architecture	27
4.5.1. Available Hardware	27
4.5.2. Interface Component	28
4.5.3. Web Application Component	28
4.5.4. Virtualization Components	29
4.5.5. Evaluation component	31
4.5.6. Storage component.....	32
4.5.6.1. Virtual machine storage.....	32
4.5.6.2. Web application database	32
4.6. FastAPI Adoption: Overcoming Flask's Shortcomings	33
4.6.1. Initial setup: Flask	33
4.6.2. Second setup: Flask + Celery	35
4.6.3. Third setup: Quart, an ASGI-compliant Flask reimplementation	36
4.6.4. Final Decision: FastAPI Migration	37

5. Implementation	39
5.1. Project structure	39
5.1.1. Technologies	39
5.1.1.1. Python	39
5.1.1.2. SQLite	40
5.1.2. app/	40
5.1.3. inventory/	41
5.1.4. logger/	42
5.1.5. nornir_lib/	43
5.1.6. proxmox_api/	44
5.1.6.1. Error Handling	44
5.1.7. gns3_api/	45
5.2. Web Application Components	46
5.2.1. Authentication Module	47
5.2.2. Exercise Management Module	47
5.2.3. VM Control Module	47
5.3. Asynchronous Processing with FastAPI	47
5.3.1. Differences Between Asyncio And Celery	48
5.4. GNS3 Customization and Configuration	49
5.5. Proxmox API Usage	50
6. Testing & Evaluation	52
6.1. Performance Evaluation	52
6.1.0.1. Task 1: Template VM Creation	53
6.1.0.2. Task 2: VM Cloning	54
6.1.0.3. Task 3: VM Deleting	54
6.1.1. I/O Problems during Batch VM Operations	54
7. Conclusion & Future Work	58
Bibliography	58

List of Figures

2.1. A simple network topology example in the GNS3 Web UI	7
4.1. A diagram showcasing how users interact with the system's resources on a high level	27
4.2. A diagram showcasing a high level overview of the system's main components	30
4.3. A diagram of our database	33
6.1. Boxplot comparison of execution time variability for Template VM creation.	53
6.2. Execution time to clone increasing number of VMs across different implementations.	54
6.3. Grouped bar chart showing VM operation times. Cloning time is labeled to highlight the rising trend. Results from Flask running purely sequential code	55
6.4. Bar chart showing VM cloning times after cleanup, at baseline, and in the worst case scenario.	57

List of Abbreviations

API Application Programming Interface. 6, 8, 17

ASGI Asynchronous Server Gateway Interface. 11, 13– 16

CS Computer Science. 1

DCC Department of Computer Science. 4

GNS3 Graphical Network Simulator-3. 6, 14, 16– 18

HTTP Hypertext Transfer Protocol. 6, 10, 14, 16– 18

IOS Internetworking Operating System. 7

IOU IOS on Unix. 7

JSON JavaScript Object Notation. 6, 15, 17

KVM Kernel-based Virtual Machine. iv, 9

LXC Linux Containers. iv, 9

OAS OpenAPI Specification. 15

Proxmox VE Proxmox Virtual Environment. 8, 9, 14, 16– 18

QEMU Quick Emulator. 7, 9

REST Representational State Transfer. 6, 8

VM Virtual Machine. 8, 9

VPCS Virtual PC Simulator. 7

WSGI Web Server Gateway Interface. 10, 11, 14

1. Introduction

This chapter contextualizes challenges, outlining the limitations of current pedagogical tools and processes while framing the necessity for an automated, scalable solution. By dissecting the shortcomings of existing platforms to manual assessment burdens—we lay the groundwork for a system capable of aiding network administration education.

1.1. Problem Statement

The digital transformation sweeping across industries has created unprecedented demand for skilled computer science professionals. While most, if not all, Computer Science (CS) specializations face growing needs, network administration remains a foundational requirement. This persistent demand reflects the networks's critical role as infrastructure supporting all digital systems.

Modern organizations require professionals who can design, configure, and troubleshoot increasingly complex network environments. Effective education must therefore bridge theoretical knowledge with practical implementation, particularly through hands-on evaluations that simulate real-world scenarios. Yet current assessment methods fail to meet these needs at scale, creating a growing gap between academic preparation and professional requirements.

1.1.1 Need for Automated Evaluation

Practical evaluations are essential to prepare students for real-world challenges, enabling them to apply theoretical knowledge and develop hands-on skills. However, available evaluation methods face significant limitations.

Creating a physical network environment for practical evaluations is sure to be costly and challenging to scale for large student populations. While emulation and virtualization technologies offer cost-effective alternatives for creating flexible practice environments, they lack built-in automated assessment capabilities. Instructors must manually review each student's network topology configuration—a process that is:

- **Time-consuming:** Manual checks grow linearly with class size.
- **Error-prone:** Human reviewers may overlook misconfigurations.

- **Inconsistent:** Subjective grading criteria lead to unfair assessments.

Automating evaluations would reduce instructor workload, freeing time for student support as well as ensure consistent, objective grading and simultaneously enable immediate feedback for learners.

1.1.2 Limitations of Current Solutions

Existing approaches to network education suffer from two critical gaps:

1. **Single-vendor focus:** Most tools (e.g., Cisco's packet tracer) are designed for specific vendor ecosystems, failing to prepare students for heterogeneous real-world networks where multi-vendor interoperability is essential.
2. **Missing evaluation component:** Some education institutions' curricula, such as our case in Department of Computer Science (DCC), forgo practical assessments entirely, which means:
 - Students complete exercises without validation
 - No measurable feedback on configuration skills
 - Graduates enter industry with possible gaps in their knowledge

1.2. Aims and Objectives

Building upon the foundational work of Santos [1] in automated network topology evaluation, this project has the following technical objectives:

1. **Develop a prototypical automated evaluation environment:**
 - Create a system for assessing network administration exercises without manual intervention
 - Support multi-vendor device configurations (Cisco, Juniper, Linux-based)
 - Validate both connectivity and configuration compliance
2. **Implement a cohesive back-end system:**
 - Transform loose components from Santos [1] into a unified system
 - Develop capabilities to coordinate between:

- Infrastructure provisioning
- Network emulation
- Evaluation automation

2. Background

The main focus of this chapter is to provide the reader with the necessary background to understand the context and technical foundations of this project. The goal of the system is to automatically evaluate network topologies by validating configurations and executing tests across various devices within a virtual network.

Achieving this requires the integration of multiple technologies, as the system must support a wide range of features to deliver an automated solution. Key concerns include not only functionality but also scalability, since multiple students may interact with the platform concurrently, each requiring an isolated working environment.

To support these requirements, this chapter introduces core concepts and tools such as virtualization, web frameworks, administration automation and task processing. These components form the foundation upon which the system is built.

2.1. Overview of Used Technologies

2.1.1 Virtualization

Virtualization is the process of creating a virtual version of physical resources, such as routers, switches, or even entire computers. In the context of this project, it is used to create virtual machines to provide students with a work environment consisting of a virtual network, itself comprised of various types of virtualized devices. This approach enhances scalability and reduces costs, as it allows multiple virtual machines to be run on a single physical machine.

Virtualization can be categorized into **emulation** and **simulation**.

- **Emulation** is the process of creating a virtual version of a physical device in software, replicating its behavior exactly—including any bugs and limitations. This is useful for various things like testing software on different platforms, running legacy software on modern hardware and even running potentially harmful software in a safe isolated environment. Emulation will be used wherever possible to provide students with a work environment that matches the real world as much as possible to best develop their network skills

- **Simulation** models the behaviour of a device, without replicating the underlying hardware or software. This results in a simpler less resource intensive model, though it may not fully capture the real device's behavior. Simulation will be used to simulate the behaviour of certain, simpler and generic, network devices and PCs.

2.1.2 GNS3

Graphical Network Simulator-3 (GNS3) is an open-source graphical network emulator software that allows the user to create complex network topologies and interact with the various devices in it. It is widely used for educational purposes and is often used in preparation for professional network certifications like the Cisco Certified Network Associate (CCNA).

GNS3 employs a simple drag and drop interface to allow users to add new devices, make links between them and even add textual annotations. The software allows users to interact with the devices by way of a console or even a GUI if the device supports it. The software also allows users to export their topologies to be shared with others, which can be useful for teachers to provide students with a pre-configured topology to work on.

Additionally, the software supports packet capturing which is essential for students to develop their debugging and troubleshooting skills. Finally it can also be interacted with via a Representational State Transfer (REST) Application Programming Interface (API) which is of particular interest for this project.

2.1.3 Architecture

The software can be employed in a variety of ways due to its architecture [2] that separates the user interfaces that it offers, namely the locally installed gns3-gui as well as the browser accessible gns3-web, from the gns3-server that runs the emulations and the controller who orchestrates everything.

2.1.3.1 Controller

The controller is integrated in the gns3-server project and is responsible for communicating with all the other components of the software. The controller is a singleton, meaning there should only be one instance of it running at any given time, and it does not support concurrent requests. It is able to control multiple compute instances if so desired, each

capable of hosting one or more emulator instances, varying depending on their complexity. The controller also exposes the REST API allowing the ability to interact with the software programatically. All communication is done over Hypertext Transfer Protocol (HTTP) in JavaScript Object Notation (JSON) format and there is support for basic HTTP authentication as well as notifications via websockets.

2.1.3.2 Compute

The compute is also integrated in the gns3-server project and controls the various emulators required to run the nodes in the topology. The list of currently supported emulators is:

- **Dynamips** - Used to emulate Cisco routers and basic switching.
- **IOS on Unix (IOU)** - Used to emulate Cisco Internetworking Operating System (IOS) devices.
- **Quick Emulator (QEMU)** - Used to emulate a wide variety of devices.
- **Virtual PC Simulator (VPCS)** - A basic program meant to simulate a basic PC.
- **VMware/VirtualBox** - Used to run virtual machines with nested virtualization support.
- **Docker** - Used to run docker containers.

2.1.3.3 GUI

The GUI is composed of two separate but with mostly identical functionality, namely the gns3-gui and the gns3-web projects. The gns3-gui project is a desktop application that is used to interact with a local or remote gns3-server instance. It is written in Python and uses the Qt framework for the graphical interface. The gns3-web is a web interface that is accessible via web browser and even though it is still in a beta stage, it has all the necessary features and stability to be used as a substitute for the gns3-gui.

2.1.4 Proxmox VE

Proxmox Virtual Environment (Proxmox VE) is an open-source platform designed for enterprise-level virtualization [3]. It is based on the Debian distribution of Linux and provides a web-based interface for managing virtual machines and containers. It is widely




Figure 2.1: A simple network topology example in the GNS3 Web UI

used in data centers and cloud environments, as it provides a scalable and reliable solution for virtualization.

Proxmox VE bundles several core services that can be interacted with via shell commands, a web interface or by using the Proxmox VE REST API. These allow the user to interact with every service provided by Proxmox VE, in a plethora of ways, depending on the user's needs, skills and preferences. The web interface is the most user-friendly way to interact with the platform, as it provides a graphical interface for managing the cluster. The shell commands provide a more direct way to interact with the platform, allowing for more complex operations to be performed and opening the doors to scripting and automation. Finally, the Proxmox VE REST API allows for programmatic and remote interaction with the platform, enabling users to create custom applications that can interact with the platform.

2.1.4.1 Virtualization Technologies

Proxmox VE supports the deployment and management of two distinct types of virtualization, namely, Kernel-based Virtual Machine (KVM) and Linux Containers (LXC).

Users can interact with these virtualized environments via NoVNC, a simple web-based VNC client or Simple Protocol for Independent Computing Environments (SPICE) which is a more feature-rich protocol that provides better performance and more features than VNC. Both of these protocols support the use of a console-based interface, as well as a full desktop graphical interface.

2.1.4.2 KVM

KVM is a virtualization solution provided by the Linux kernel. It leverages the hardware virtualization extensions of modern processors to provide a full virtualization experience at near-native speeds. Supports a wide range of guest operating systems making it a good choice for general purpose virtualization.

In Proxmox VE, KVM is used as the core component for running virtual machines and is used alongside QEMU.

2.1.4.3 LXC

Containerization is an operating system-level virtualization method that packages an application and its dependencies together into an isolated environment. Contrary to traditional Virtual Machine (VM) solutions, containers don't emulate hardware or require a guest operating system relying instead on the host's kernel. This approach leads to a faster and more lightweight virtualization solution, as they consume less memory and CPU resources.

LXC creates full system containers, capable of simulating a complete Linux distribution providing users with an environment that behaves like a traditional VM but with the speed and efficiency of a container. LXC starts much faster than VMs making them ideal for scenarios requiring rapid deployment and/or scaling.

However, it's important to note that while containers offer a degree of isolation, they do not provide the same level of security as VMs. This means that while they may not always be a suitable replacement for VMs.

2.1.5 LDAP

Lightweight Directory Access Protocol (LDAP) is the foundation of user and device management in many institutions. Universities can rely on openLDAP and/or Microsoft's Active Directory (its enterprise implementation) to handle student, faculty accounts and lab computer access, amongst other things.

One of LDAP's most popular implementations, OpenLDAP, had its initial release in 1998. The protocol's longevity stems from its efficiency at handling large-scale authentication so much so that despite newer alternatives existing, LDAP remains entrenched in academic

environments due to its reliability and universal adoption. For our project, LDAP integration enables students access to the system using their existing university credentials.

Two key factors make LDAP particularly valuable for this project: its standardized approach to user management and pre-existing deployment in our target educational environments. Given the extensive use, it's desirable for our system to have the capability to interact with LDAP in order to correctly authenticate users.

2.2. Virtualized Lab Environments

The combined use of Proxmox VE as a virtualization platform and GNS3 for network emulation presents a cost-effective solution for scalable networking education. This approach offers significant benefits over physical lab infrastructures:

- **Resource Efficiency:** Single physical host can support multiple concurrent student environments
- **Operational Characteristics:**
 - Accelerated environment provisioning through templates
 - State preservation via Proxmox VE's snapshot/restore functionality
 - Support for diverse network operating systems through virtualization technologies

2.3. Python Web Frameworks for API-Based Systems

2.3.1 Python

Python is a high-level, interpreted programming language renowned for its readability and versatility. It supports multiple programming paradigms, including procedural, object-oriented, and functional programming, making it suitable for a wide array of applications. In the context of this project, Python serves as the primary programming language as its extensive standard library and supportive community contribute to efficient development and maintenance of the project's codebase.

2.3.2 WSGI

The Web Server Gateway Interface (WSGI) is a pivotal standard for Python web application deployment, defining a consistent interface between web servers and Python web applications/frameworks.

Prior to WSGI's introduction[4], Python web frameworks were typically written against various server-specific APIs such as CGI, FastCGI, or mod_python. This diversity led to compatibility issues, limiting developers' choices of web servers and frameworks, as not all frameworks supported all web servers and vice-versa. To address this fragmentation, WSGI was created as a standardized interface, promoting portability and flexibility in deploying Python web applications.

WSGI serves as a bridge, enabling web servers to communicate with Python applications. It specifies a simple and universal interface for web servers to forward requests to Python applications and for those applications to return responses. This standardization allows developers to choose from a variety of web servers and Python frameworks without compatibility concerns.

Introduced in 2003 as PEP 333, WSGI was later updated to PEP 3333 in 2010 to accommodate Python 3. These specifications outline how web servers and Python applications should interact, ensuring a consistent and reliable deployment environment across different platforms.

The WSGI standard consists of two main components:

- **Server/Gateway Side** - Responsible for receiving HTTP requests from clients and passing them to the Python application. Then receives the response from the application and forwards it to the client.
- **Application** - The Python application that processes requests and returns responses.

Additionally WSGI has support for middleware components. WSGI middleware is a Python callable that wraps another WSGI application to observe or modify its behavior. Middleware can perform various functions, including request preprocessing, response postprocessing, session management, and security checks. This modularity allows developers to add functionality to their applications in a reusable and maintainable manner.

The separation defined by WSGI allows for flexibility and scalability in deploying Python web applications.

Python WSGI applications often use built-in servers, during development, provided by frameworks like Flask. However, these servers typically aren't fully featured and aren't suitable for production environments. In production, WSGI servers act as intermediaries

between web servers (e.g., NGINX or Apache) and Python applications, handling incoming requests and serving responses efficiently.

2.3.2.1 Flask

Flask is a web application micro framework written in Python, adhering to the WSGI standard, designed to facilitate the development of web applications by providing essential tools and features. Classified as a microframework, Flask does not require particular tools or libraries, instead choosing to focus on simplicity and extensibility[5].

An example of how easy it is to develop a basic web application with flask is provided in the following small piece of code.

Algorithm 1 Flask Hello World

```

1: from flask import Flask
2: app = Flask(__name__)
3:
4: @app.route('/')
5: def hello_world():
6:     return 'Hello, World!'
7:
8: if __name__ == '__main__':
9:     app.run()

```

2.3.3 ASGI

Asynchronous Server Gateway Interface (ASGI) is an interface specification for Python web servers and applications. It is considered a spiritual successor to WSGI, designed to provide a standard interface for asynchronous communication. ASGI was developed to address the limitations of WSGI, which was primarily designed for synchronous applications. Unlike WSGI, ASGI supports handling multiple requests concurrently, making it suitable for modern web applications that require real-time features such as WebSockets, long-lived connections, background tasks or the use of Python's async features.

As development progressed, asynchronous task handling became a more central requirement, initially addressed by integrating task queues. However, due to resource overhead and deployment complexity, they were phased out. This shift prompted an evaluation of frameworks that offered native support for asynchronous operations.

2.3.3.1 FastAPI

FastAPI is a modern, high-performance web framework adopting the ASGI standard. It leverages open standards, such as OpenAPI Specification (OAS), for defining path operations, parameters, and more, which in turn is based on the JSON schema. FastAPI relies entirely on Python type declarations, making it more intuitive and lowering the barrier to entry to new developers. This approach also simplifies the understanding and maintenance of the codebase.

Built on top of Starlette, a lightweight ASGI framework, and Pydantic, a data validation library. FastAPI combines the strengths of both to provide a powerful and flexible framework for building APIs with automatic data validation, serialization and documentation generation, all of which significantly enhance developer productivity.

Another key feature of FastAPI, being ASGI-compliant, is its built-in support for asynchronous programming, allowing developers to write non-blocking code using Python's *async/await* keywords. This is particularly useful for I/O-bound operations, such as database queries or network requests, as it allows the application to handle multiple requests concurrently without blocking the application which is essential in projects such as this one where multiple concurrent HTTP calls are made to interact with multiple devices and services concurrently, such as GNS3 and Proxmox VE.

Another powerful feature of FastAPI is its dependency injection system, that is very easy to use as it is automatically handled by the framework itself. This allows for a clean and modular codebase, as dependencies can be easily injected into the various components of the application. This is especially useful in larger applications, where managing dependencies can become complex and cumbersome.

A change from Flask to FastAPI laid the groundwork for more efficient handling of I/O-bound operations—such as network interactions with Proxmox VE or GNS3, which will be of importance in future iterations of the project while also streamlining development thanks to FastAPI's built-in request parsing, background task support, and integrated dependency injection system.

2.4. Long running task processing approaches

2.4.1 The need for asynchronous processing in API-heavy applications.

Modern API-driven applications benefit tremendously from asynchronous programming paradigms to handle concurrent operations efficiently. Traditional synchronous execution models, where each request blocks thread execution until completion, prove inadequate for systems requiring high throughput and responsiveness. This limitation becomes particularly apparent in projects like ours, which relies heavily on HTTP calls to various devices and services.

The asynchronous model, implemented through Python's `async/await` syntax, offers several critical advantages:

- **Improved Resource Utilization:** A single thread can manage multiple concurrent I/O operations by yielding control during waiting periods
- **Enhanced Scalability:** Systems can handle higher concurrent user counts with the same hardware resources
- **Responsive Performance:** The application remains reactive even during long-running operations

Although Python has native asynchronous capabilities, libraries must be written with these in mind, meaning some may have limited or even no support for these capabilities.

2.4.2 Asyncio

Asyncio is Python's built-in library for writing concurrent asynchronous code. It serves as the foundation for asynchronous operations in many Python frameworks, enabling high-performance networking, web servers, database connections, distributed task queues, etc..

Asynchronous I/O can be useful in cases of time-consuming operations, as while awaiting the finish of said tasks, it relinquishes control so that other code can run in the meantime. This approach is particularly well-suited for I/O-bound operations—such as network communication, file access, or database queries—where tasks would otherwise spend a significant amount of time waiting for external operations that are outside of our control to complete. Rather than blocking the entire application during such waits, Asyncio allows

other tasks to execute in the meantime, leading to more efficient resource utilization and improved throughput.

Overall, Asyncio provides the concurrency model that underpins efficient I/O performance. By embracing this model, the project benefits from improved responsiveness, lower latency, and better scalability—especially under workloads that involve heavy interaction with external services.

Frameworks that leverage these capabilities natively, provide the foundation for building responsive, scalable API services.

2.4.3 Celery

Celery is an open source distributed task queue focused on real-time processing but also offers support for task scheduling. It is implemented in Python, but the underlying protocol can be implemented in any language. Celery requires a message broker to function, such as Redis or RabbitMQ, which are responsible for queuing and distributing tasks from producers (clients) to consumers (workers).

When integrating Celery into their projects developers must mark functions they want to be processed as tasks with Celery provided decorators (e.g. `@app.task`) which allows workers to then execute them asynchronously. The system shines in projects requiring heavy computational or scheduled jobs, but brings with it non-negligible operational, developmental and resouce overhead, doubly so if the project didn't already include the use of message brokers.

2.5. System administration automation tools

Modern system administration increasingly relies on automation tools to manage complex infrastructure while maintaining reliability and reproducibility. In our context, configuration management systems serve as the foundational layer for ensuring well-configured device states across network environments.

2.5.1 Nornir

Nornir is an open-source automation framework written in Python, designed to provide a flexible and efficient approach to network automation tasks[6]. Unlike other automation tools that utilize customized pseudo-languages, Nornir leverages pure Python code, offering developers the full power and versatility of the Python ecosystem.

Nornir supports multi-threaded task execution, allowing operations to run parallel across multiple devices. This capability enhances efficiency and reduces the time required enabling easy scaling to a large number of devices.

The framework provides a robust inventory management system, enabling the organization of devices into groups and the assignment of specific tasks to these groups. This structure facilitates targeted automation and simplifies complex network operations.

Finally, thanks to Nornir's architecture, it is highly extensible through its plugin system, allowing users to create custom plugins for inventory management, task execution, and result processing. This modularity ensures that Nornir can adapt to a wide range of network automation scenarios.

Nornir makes it easy to write reusable tasks for configuration management and state validation which makes it highly desirable in the context of this project. Its ability to handle concurrent operations will also ensure it can scale alongside the rest of the project.

2.5.2 Ansible

Ansible is a widely adopted open-source automation platform that simplifies configuration management, application deployment, and task automation through a declarative YAML-based approach. Unlike imperative scripting solutions, Ansible employs playbooks to define system states, making automation accessible to both developers and operations teams while maintaining robust capabilities for complex workflows.

The platform operates on an agentless architecture, utilizing Secure Shell (SSH) for connectivity, which eliminates the need for persistent software on managed nodes. This design choice significantly reduces deployment overhead while maintaining secure communication through standard protocols. Ansible's push-based execution model allows for immediate task execution across entire device inventories without requiring pre-installed clients.

Ansible remains a valuable tool in task automation and orchestration but, as was already discussed in[1] there are several barriers to the adoption of Ansible in this project, mainly the difficulties encountered by utilizing the Telnet protocol for communications with network devices that dont support the SSH protocol .

3. Related Work

This chapter focuses on placing the current project within the context of existing solutions and related work. The primary goal of this project is to develop a system capable of automatically evaluating network topologies by validating device configurations and executing tests across a virtual network.

While automated assessment systems are well established in the field of programming education—receiving student-submitted code and running it against predefined test cases—equivalent systems for network exercises are far less common. Tools like Mooshak and similar platforms have proven effective for evaluating programming assignments and are widely adopted in academic settings.

At first glance, adapting these approaches to network topologies might seem straightforward. However, network evaluation introduces unique challenges such as the need for per-student virtual environments, real-time communication with multiple devices, and stateful, distributed configurations. This chapter explores existing tools like Mooshak and Packet Tracer, highlighting their capabilities, limitations, and how this project builds upon or diverges from them.

3.1. Programming Evaluation Systems

While not directly related, they are the main inspiration for this project. Programming evaluation systems are widely deployed in universities and other educational institutions. These systems receive, as input, code from students and subsequently run tests on it, outputting a score and even being configurable to provide students the first test case that they failed in, guiding students to the solution without handing it out.

The main differentiator between these systems and the one proposed in this project is the ability to solve a network exercise using multiple configurations across multiple devices, while programming evaluation systems will expect the same output every time, given the same input.

Another key difference is the fact that programming evaluation systems dont always provide a working environment for the students to test their code, owing to the fact that students might prefer to user their own development environment for initial development and testing. This project aims to provide a working environment for students, as setting

up a networking lab can be a daunting task for students, especially when they are just starting out. By providing a pre-configured environment, students can focus on learning the concepts and skills they need to succeed in their studies, rather than spending time troubleshooting their setup.

3.1.1 Mooshak and lessons learned

In our context, in the DCC, Mooshak is commonly deployed to be used in the context of classes, exams and even programming contests.

Mooshak is a web-based system for managing programming contests and also to act as an automatic judge of programming contests [7]. It supports a variety of programming languages like Java, C, etc. Under each contest students will find one more problem definitions each containing varying sets of test cases in input-output pairs. After submitting their solution, the system will compile and run the code against the test cases giving a score based on the the amount of test cases passed.

Mooshak provides a structured approach to test coding and problem solving skills. It begins by offering a problem statement coupled with an optional image and an example test case, normally in the form of input and expected output. Users can submit their proposed solution by uploading a file with their code. The system then evaluates the provided solution against multiple pre-defined test cases, validating the output against the know-good output, and giving feedback in the form of a score based on the number of test cases passed. The system may also be configured to have time and/or memory constraints, to ensure that temporal and spatial complexity are also taken into account.

All of these, serve to provide a thorough evaluation of the student's solution, which can help guide a student to better their coding and problem solving skills.

The system can also differentiate between differing types of errors, such as not giving the expected output, poorly formatted output, failure to compile or even exceeding the time limits. Mooshak also includes some features designed to drive competition between students, like a real time leaderboard and the ability to have more than 100% of the score for a given contest.

The system however is not without its limitations as it uses plain text files for its test cases and validates the output of student's code character by character, which can lead to false negatives if the output is not formatted exactly as expected.

3.2. Cisco Packet Tracer

Cisco Packet Tracer is a network **simulation** tool developed by Cisco Systems, widely used in academic environments to teach networking concepts and prepare students for certifications such as the Cisco Certified Network Associate (CCNA). It offers a visual interface for building and simulating virtual network topologies using a variety of Cisco devices, including routers, switches, and end devices.

While Packet Tracer is highly accessible and effective for introducing networking fundamentals, it is a closed-source, proprietary tool limited to simulating Cisco hardware and IOS features. Its functionality is optimized for teaching purposes rather than for flexibility, extensibility, or integration into larger automated workflows.

In contrast, this project aim to allow for a more realistic and extensible lab environment. The use of real operating systems and support of a wide range of vendor platforms for routers and switches, aswell as Linux-based virtual machines is highly desirable. This allows for a more realistic experience, as students will be able to work with the same tools and operating systems that they will encounter in real-world scenarios.

Therefore, while Cisco Packet Tracer remains a valuable educational tool, the needs of this project called for a more flexible and open architecture.

4. System Architecture & Design

This chapter outlines the architecture of the proposed system, detailing the key components and how they interact to enable evaluation of student-submitted network exercises.

The system is designed to provide each student with a working environment where custom virtual network topologies can be deployed, configured, and tested. To achieve this, the platform integrates several technologies—such as GNS3 for network emulation, Proxmox VE for virtualization, and Nornir for configuration testing—alongside an asynchronous web-based API layer for user interaction and system communications.

This section provides a high-level overview of the system, the rationale behind its design choices, and the fundamental components that make up its architecture.

4.1. Functional Use Cases

This section outlines the key functional use cases that the system must support to fulfill its intended requirements. Functional use cases describe the interactions between the users (actors) and the system, detailing the expected behavior and workflow for each core functionality. These use cases serve as a foundation for understanding the system's design and implementation.

4.1.1 System Supports Both Directory Integration and Local Authentication

Actor: System Administrator

Main Flow:

1. The administrator configures the authentication method.
2. If directory integration is enabled, LDAP or similar service becomes responsible for credential validation.
3. If local authentication is selected, user credentials are stored and verified locally.
4. During login, the system checks which method is active and authenticates accordingly.

Postconditions: Users can authenticate via the configured method: directory service or local credentials.

4.1.2 Teacher Creates a New Exercise

Actor: Teacher

Preconditions: The teacher is authenticated in a priviledged account.

Main Flow:

1. The teacher fills out a exercise creation form with the exercise metadata.
2. The teacher enters a template VM ID into the form.
3. The teacher submits the form.
4. The system clones the selected template.
5. The system uploads the GNS3 project file provided in the form to the newly cloned VM.
6. (Optional) The system performs form-defined pre-configurations to the GNS3 topology.
7. The system transforms the new VM into a template ready to be cloned for use by students.

Postconditions: The exercise and related specific VM are created and stored in the system and become visible to enrolled users.

4.1.3 System Supports Enrollment of Users into Exercises

Actor: Teacher

Preconditions: Atleast one exercise exists in the system.

Main Flow:

1. The teacher selects an exercise.
2. The teacher navigates to an enrollment interface.
3. The teacher uploads a list of new users.
4. (Optional) The system checks the list of users against existing directory services.
5. If necessary, new users are registered automatically.

6. The system validates if the list contains already enrolled users.
7. The system creates VMs for non-enrolled users.
8. The system creates associations between the users, the VMs and the exercise.

4.1.4 Student Starts Their Exercise VM

Actor: Student

Preconditions: Student is enrolled in selected exercise.

Main Flow:

1. The student logs in and navigates to the assigned exercise.
2. The student clicks the `Start VM` button.
3. The system sends an API request to the Proxmox server to start the VM..

Postconditions: The VM is started and its GNS3 instance becomes accessible to the student.

4.1.5 Student Submits Solution

Actor: Student

Preconditions: The student's VM and its GNS3 instance are running.

Main Flow:

1. The student clicks the `Submit` button for the exercise.
2. The system (via Nornir) connects to the student's VM.
3. Commands are executed and their output is validated.

Postconditions: The validation results are saved and feedback is made available.

4.1.6 Teacher Reviews Student Results

Actor: Teacher

Preconditions: At least one student has submitted their solution.

Main Flow:

1. The teacher logs into the system and selects an exercise.
2. The system displays a list of students along with their submission statuses.
3. The teacher selects a student to view detailed feedback and logs.

Postconditions: The teacher reviews the validation output for grading or feedback.

4.1.7 Student Views Exercise Feedback

Actor: Student

Preconditions: The student has at least one completed submission.

Main Flow:

1. The student logs into the system and navigates to submission history.
2. The student selects a submission to view its result.
3. The system displays validation output and feedback.

Postconditions: The student sees the outcome of the submission and understands any errors.

4.2. System Architecture Overview

The architecture is divided into several key components, each responsible for a specific aspect of the system's functionality. The main components of the system architecture are as follows:

- **Web Application:** The web application serves as the main interface for users to interact with the system. It provides endpoints for, amongst others, evaluation, creation and viewing available exercises. The application is designed to be asynchronous where possible, allowing for efficient handling of multiple requests simultaneously.
- **Proxmox VE:** Proxmox VE is responsible for creating and managing VMs that host the network devices used in the exercises. This component interacts with the web application and all communication is done asynchronously through the Proxmox VE REST API, which allows for efficient communication, keeping the web application responsive, while also keeping the components decoupled.

- **GNS3:** GNS3 is used to emulate all the components of the virtual networks to be configured by students, using various types of virtualization detailed earlier. Communication with GNS3 is done through the GNS3 REST API by the web application
- **Nornir:** This automation framework is used for validating device configurations. It connects to the virtualized devices, executes commands, and compares the output to expected results to determine correctness. Currently this component is integrated into the web application

4.3. Proxmox VE

Proxmox VE functions as the virtualization backbone of the system, enabling the creation and management of Linux-based VMs which in turn host services for use by students. Each student's VM runs a lightweight Linux-based operating system with a dedicated GNS3 instance, providing a self-contained environment for deploying and configuring virtual networks and their components.

All Proxmox VE-related operations—such as cloning, starting, templating, and deletion—are fully automated and triggered by the web application. Under normal operating conditions, no manual intervention via the Proxmox VE web interface or shell utilities is required once the initial setup is complete. Such intervention should only be necessary when the system's error handling mechanisms aren't capable of automatically resolving errors. To securely execute these operations, the application authenticates to the Proxmox VE API using token-based authentication, ensuring that only authorized and properly configured processes can interact with the Proxmox VE infrastructure.

4.3.1 Why Proxmox VE?

Proxmox VE was chosen for several compelling reasons that make it ideal for it to be chosen as our virtualization platform. First, it's completely free to use for all core functionality, with no hidden costs or licensing traps. Unlike proprietary solutions that charge per CPU core or socket, Proxmox VE lets us scale up our infrastructure without worrying about licensing fees.

The platform's support for both containers and VMs within a single management interface gives us tremendous flexibility. We can run lightweight LXC for applications that don't require a full VM while using full VMs where required seamlessly. This hybrid approach would not be as straightforward with other solutions, like VMware ESXi.

We also value storage system's flexibility with LVM-thin provisioning allows efficient snapshotting of student environments while maintaining good performance.

Looking ahead, Proxmox VE's built-in support for technologies like software-defined networking and a robust role-based access control system means our project still has room to grow into Proxmox VE. The active open-source development community ensures continuous improvements without vendor lock-in.

4.3.2 Proxmox VE Limitations

During development, we encountered several challenges when interfacing programmatically with Proxmox VE. One of the most significant issues stemmed from the platform's limited visibility into non-instantaneous operations - particularly for tasks like VM cloning, where the system did not provide task ids in the HTTP responses. This forced us to implement custom polling mechanisms to reliably determine operation completion states where possible.

A more critical limitation emerged in Proxmox VE REST API's resilience characteristics. During stress testing, we discovered that even moderate request volumes done using a single machine running sequential code could overwhelm the single-node cluster's management daemon, triggering frequent HTTP 500 errors. These reliability constraints necessitated the development of protective measures including exponential backoff retry logic and strict web application concurrent request limiting when interacting with Proxmox VE to maintain system stability.

One final limitation discovered during performance testing of batch creation and deletion of VMs was that, on occasion, their disks would fail to be removed. This issue could not be detected through HTTP responses alone and was only observable by inspecting the Proxmox VE logs. As more disks failed to be removed, they accumulated as unused entries in storage, eventually leading to a noticeable performance regression. This specific limitation will be discussed in more detail in a later chapter.

4.3.3 Proxmox VE Firewall

Proxmox VE comes bundled with an iptables-based firewall implementation that can be enabled and configured at different levels.

The Proxmox host-level firewall provides essential features for securing student work environments, during examination periods, preventing student machines and the virtualized network equipments in them from communicating with outside networks.

This is done by adding firewall rules at the host level, meaning to each relevant student VM, that disable communications in both directions, with the exception of the machine that is responsible for configuration validation and the machine the student is working from.

By default this behavior is not active and must be enabled on an as-needed basis, such as when a controlled assessment environment is required for more rigorous situations such as examinations.

In future iterations it may also be valuable to develop this further and making this feature less rigid as it may be interesting to have exercises that communicate with devices on the internet.

4.3.4 Maximizing resource usage

During development, we attempted to minimize resource usage where possible to achieve higher levels of scaling potencial. The introduction of the GNS3 web interface removed the need for students to interact with the VM directly, allowing for operation in a headless manner. This eliminated the need for a desktop environment reducing memory overhead which improved scalability.

4.3.4.1 Exploration of containers as a full substitute for VMs

We further explored replacing VMs entirely with containers, which promised additional resource savings. However, this approach proved unfeasible due to fundamental technical constraints. Effective emulation requires KVM acceleration, which presents two problematic scenarios in containerized environments: either running unprivileged containers without KVM access, resulting in unacceptable performance degradation or configuring privileged containers or containers with KVM passthrough, introducing serious security vulnerabilities.

Given that software emulation without KVM acceleration delivers poor performance for interactive use, we abandoned containerization as a complete VM replacement for GNS3 hosts. The VM-based architecture remains necessary to maintain both performance through hardware acceleration and proper isolation between student environments.

4.3.5 VM Lifecycle

The lifecycle of a VM begins when a new exercise is created by a priviledged user through the web application. Upon exercise creation, the platform automatically clones a pre-configured base template VM stored in Proxmox VE. This new instance undergoes a configuration process where the provided GNS3 project file is imported and a series of user-defined commands are executed across the provided network topology. Once the setup is finalized, the configured VM is converted into a new template VM that is tailored to that exercise.

When students are enrolled in an exercise, the system generates individual work environments, by creating linked clones from these exercise-specific templates. Each student receives their own isolated VM instance that precisely mirrors the original template's configuration. This cloning approach ensures both consistency across student environments and rapid provisioning, as linked clones avoid the overhead of full disk copies while maintaining the template's baseline configuration. The use of linked clones significantly reduces both storage requirements and deployment time compared to traditional full cloning methods.

4.4. GNS3

GNS3 serves as the core network virtualization component in our system, providing the capability to emulate various network devices and topologies. The platform was selected for several key advantages: its remote web-based interaction, the intuitive drag-and-drop interface simplifies usage, and its broad device support accommodates both terminal-based and GUI network equipment and full computers. Additionally, GNS3's API allows for programmatic interaction, which proves essential for automation within our environment.

Currently, the system requires manual preparation of the base GNS3 template VM. During initial setup, an administrator must first create and configure a new VM, then proceed to install and set up the GNS3 environment. The final preparation step involves importing all necessary device images, including routers, switches, and other equipment that will be available for student exercises.

This base template then serves as the source for all subsequent student instances through Proxmox VE's cloning functionality. While this manual setup process adds initial configuration overhead, it ensures complete control over the base environment and allows for

careful curation of the included device images.

The system achieves scalability through multiple VM instances running in Proxmox VE, each hosting an independent GNS3 environment. This architecture enables concurrent usage by multiple students on a single physical host. For future expansion the system can be installed on a machine with a higher core count and RAM amount and additionally be expanded horizontally by adding additional nodes to the Proxmox VE cluster, allowing the platform to accommodate growing numbers of users without requiring complete architectural changes.




Figure 4.1: A diagram showcasing how users interact with the system's resources on a high level

4.5. High-level architecture

4.5.1 Available Hardware

The current deployment hosts all internal system components (those shown in the architecture diagram excluding the external LDAP instance) on a single physical server with the following specifications.

Table 4.1: System Hardware Specifications

Component	Specification
Processor	Intel Core i7-9700K
Memory	32GB DDR4 @ 2666MHz
Storage	1TB Samsung 970 EVO Plus NVMe SSD
Graphics	NVIDIA GTX 1650

This machine’s specifications, while capable enough for development purposes, create inherent memory constraints. With 32GB of available RAM, practical VM allocation becomes the primary bottleneck. For instance, when deploying GNS3 instances each configured with 4GB of memory, the system can maintain only seven active VMs simultaneously. This limitation accounts for Proxmox VE’s own memory overhead before inducing SWAP file usage, which would degrade performance.

4.5.2 Interface Component

This component is comprised by two distinct UIs accessible via standard browsers. For administrative functions and exercise management, users interact with Jinja2-rendered HTML pages delivered by the web application. These handle all developed features, like user authentication, VM interaction etc.

When working on networking exercises, users transition to the GNS3-web interface. This dedicated environment provides direct access to the user’s virtual network devices, as required by each exercise scenario. This integration ensures users experience a cohesive workflow from exercise selection to practical implementation without needing multiple authentication steps or application switches.

4.5.3 Web Application Component

The web application serves as the primary interface through which users interact with the system. It is built using the FastAPI framework, following a migration from an earlier prototype developed using Flask, and follows an asynchronous-first, modular architecture that provides scalable interactions with other system components.

The application exposes a REST API that supports endpoints for user authentication, exercise creation, virtual machine management, and configuration validation. It acts as the coordinator for the entire system, triggering operations in Proxmox VE, GNS3, and Nornir based on user actions.

Wherever possible, asynchronous I/O is employed to prevent blocking during operations such as API calls to Proxmox VE. Multiprocessing is also utilized to handle configuration validation. This keeps the system responsive and performant, especially when handling multiple simultaneous requests from different users.

Internally, the application is designed to be stateless and maintain minimal runtime state. Most essential information—such as user accounts, defined exercises, and student-to-VM mappings—is persisted in a relational database rather than stored in memory. Configuration parameters such as the cryptographic secret key, the Proxmox VE host IP address, database connection, as well as LDAP-specific realms and distinguished names, are injected through environment variables. This approach ensures that deployment-specific settings are decoupled from the application logic, promoting flexibility, portability, and security.

To ensure maintainability and modularity, interactions with external services like Proxmox VE and GNS3 are isolated in dedicated modules. These serve as abstraction layers between the application logic and third-party APIs, exposing clean, reusable interfaces while hiding low-level implementation details. For example, Proxmox VE-related operations such as VM creation and deletion are handled in a separate module (e.g. `app/services/proxmox.py`), which themselves interface with lower level implementations (`proxmox_api/proxmox_vm_actions.py`) as are all GNS3 – related tasks. This separation of concerns improves the structure.

To help with development and testing, the application automatically generates OpenAPI-compliant documentation, a feature of FastAPI, allowing developers to explore and interact with available endpoints. This self-documenting behavior streamlines integration testing and encourages a more agile development process.

Finally, to safeguard user data and infrastructure control points, the application enforces secure authentication mechanisms using JSON Web Token (JWT) ensuring that only authorized users can trigger actions on shared resources.

4.5.4 Virtualization Components

The system employs a hybrid virtualization approach using Proxmox VE as the foundational platform. The usage of containers was explored but it was found unsuitable for our main use case of virtualization, GNS3 instances. However there remains one valid usage for containers for the project, which is hosting the web application. However this component may also be optionally hosted in a separate physical machine.

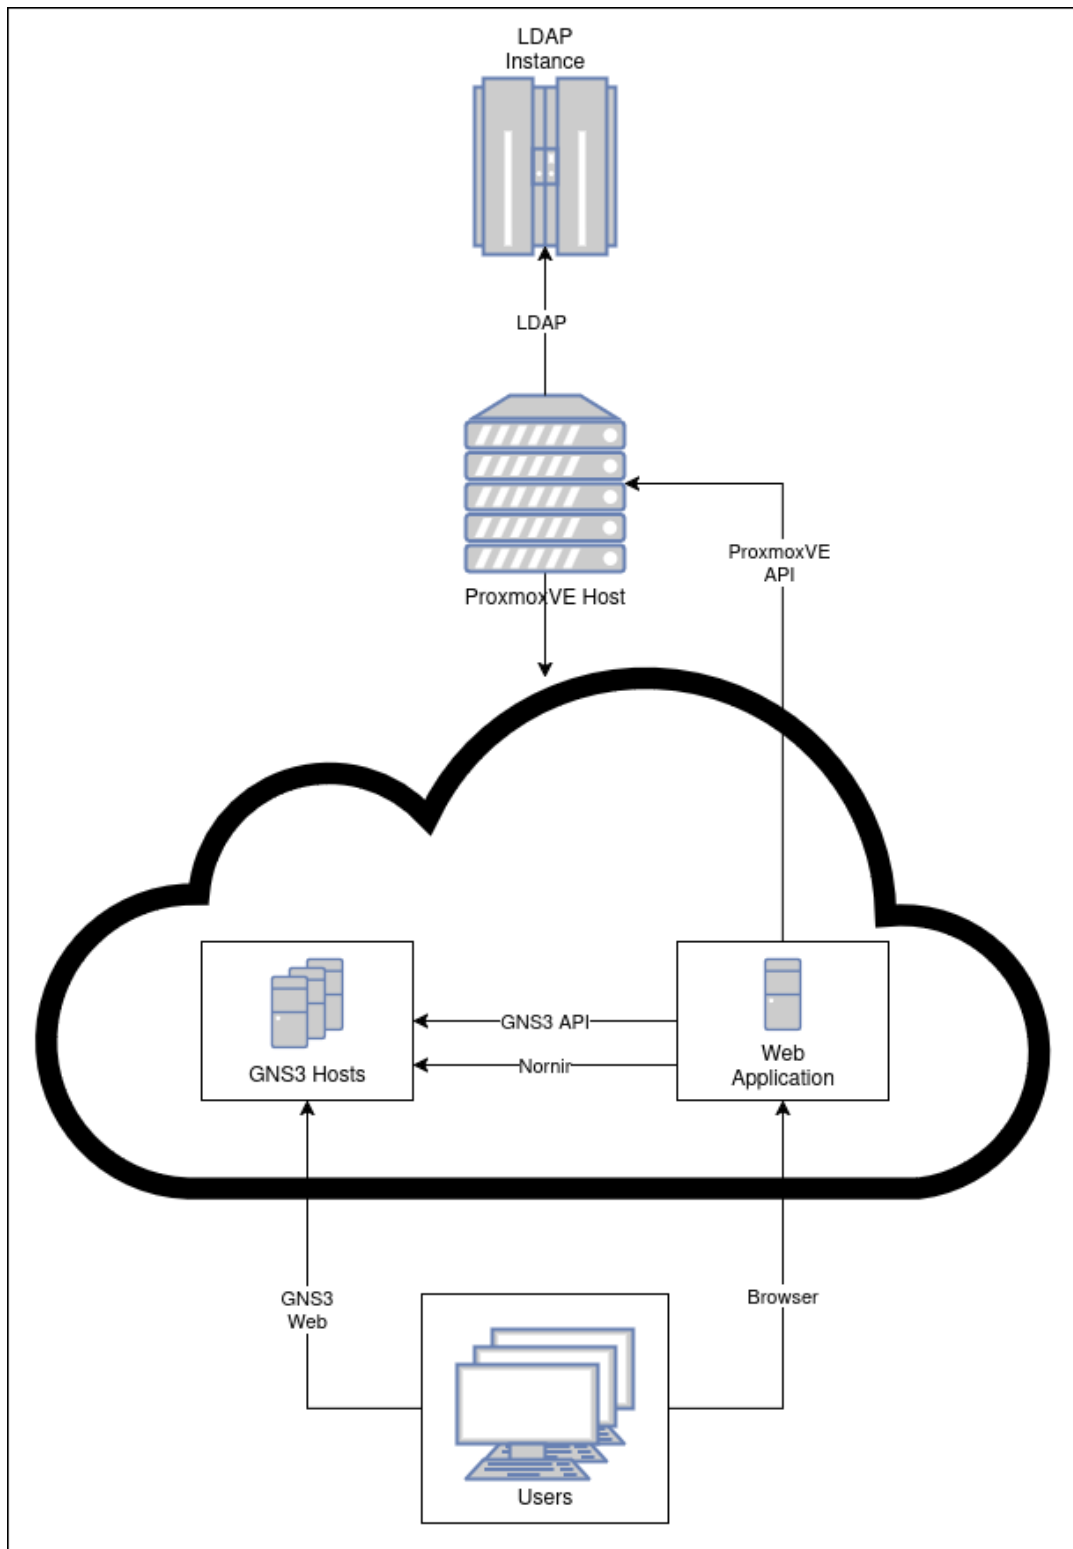


Figure 4.2: A diagram showcasing a high level overview of the system's main components

For network emulation, the system utilizes full KVM-based virtual machines, each hosting a GNS3 instance. These VMs provide the necessary hardware virtualization support for nested device emulation, particularly crucial for fast virtualization. Finally, by the use of linked clones and storage-efficient backing filesystems, allows the system to rapidly provision VMs while minimizing storage usage.

4.5.5 Evaluation component

The system employs a modular evaluation framework built on Nornir to validate configurations across virtualized network devices. At its core, this component utilizes specialized Python classes called "modules" that encapsulate platform-specific validation logic. Each module is responsible for three key functions: identifying the target device's platform (such as Cisco IOS, Linux, or VPCS), executing the appropriate validation commands for that platform, and interpreting the command output using regular expressions to determine configuration correctness.

The architecture follows an object-oriented design pattern with a base `CommandModule` class that handles common functionality. This parent class manages the Nornir inventory initialization and provides essential methods like platform detection and command execution. The actual validation logic is implemented in child classes that inherit from `CommandModule`, with each subclass specializing in a particular type of network test. For example, the included `PingModule` implements platform-specific variants of the ping command and corresponding response interpretation methods. This design promotes code reuse while allowing easy extension for new test types, as developers can create additional modules by simply extending the base class and implementing the required platform-specific methods.

Configuration validation occurs through a multi-stage process. When a test is initiated, the system first identifies the target device's platform through Nornir's inventory system. It then dispatches the appropriate platform-specific command variant, such as the Cisco IOS-style ping command for cisco routers versus the Linux `ping -c` syntax for Linux hosts. The module captures and sanitizes the raw command output, removing terminal control sequences and other artifacts before applying regular expressions to assess the results. For connectivity tests like ping, the interpretation logic calculates success rates against a configurable tolerance threshold defined in the system constants.

The evaluation framework supports several advanced features to enhance reliability and debugging. Command timeouts are managed to prevent hanging operations, with a default window that can be tuned as needed. Future extensions could incorporate snapshot functionality, allowing the system to capture and compare device states at different points during an exercise, though this capability is not currently implemented in the base version. The modular architecture ensures such enhancements can be added without disrupting existing validation workflows.

4.5.6 Storage component

The system has two main components regarding storage, one for the database needs of the web application and one for the disks of the VMs.

4.5.6.1 Virtual machine storage

LVM Thin Provisioning (LVM-Thin) is an efficient solution for creating and managing virtual machines (VMs) by optimizing storage usage and improving performance. Unlike traditional Logical Volume Manager (LVM), which pre-allocates disk space, LVM-Thin allows dynamic allocation, meaning storage is consumed only as the VM writes data—ideal for environments like ours where multiple VMs share the same storage pool. When combined with Copy-on-Write (CoW) snapshots, LVM-Thin enables rapid VM cloning and backup operations. For instance, a base VM image can serve as a template, and new VMs are created as linked clones that initially share all data blocks with the original. Only when a VM modifies its disk does LVM-Thin allocate new blocks, significantly reducing storage overhead. This approach not only saves disk space but also speeds up VM deployment, making it a great choice for our project. Additionally, since snapshots are space-efficient, in the future, we can maintain multiple VM checkpoints without worrying about excessive storage consumption—as long as the thin pool is monitored to avoid overprovisioning. Overall, LVM-Thin provides a scalable, high-performance storage layer for virtualization with minimal waste.

4.5.6.2 Web application database

The SQLite database serves as the central repository for all application data, leveraging SQLAlchemy as an ORM layer that combines Pydantic validation with SQLAlchemy's

database capabilities. This hybrid approach provides both runtime type safety and efficient database operations, while Alembic handles schema migrations to accommodate evolving data requirements.

The database schema organizes information across several interrelated models. User management builds upon a base `CustomBase` class that automatically tracks creation timestamps, with the `User` model storing authentication credentials, administrative privileges, and relationships to both submissions and virtual machine instances. The `Exercise` model captures lab configuration details, including JSON-serialized validation rules and device configurations stored as text fields due to SQLite's native type limitations. VM provisioning is managed through the `TemplateVm` and `WorkVm` hierarchy, where template instances maintain the base GNS3 project configurations and spawned work environments link back to both users and exercises. The `Submission` model completes the core data structure by tracking student attempts, scores, and evaluation outputs while maintaining referential integrity through SQLAlchemy relationships.

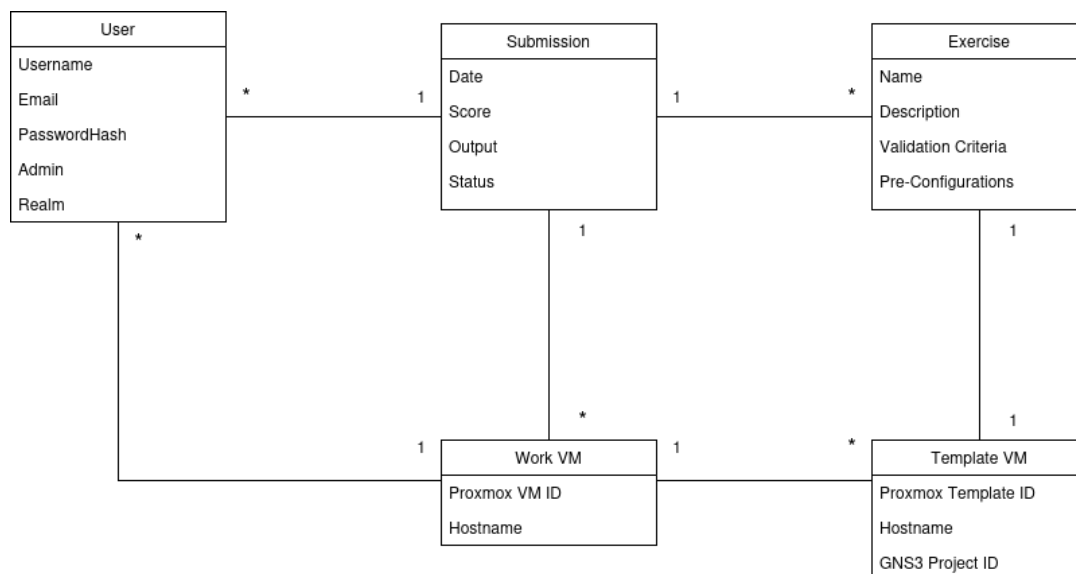


Figure 4.3: A diagram of our database

4.6. FastAPI Adoption: Overcoming Flask's Shortcomings

4.6.1 Initial setup: Flask

Initially, Flask served as the framework for the web application, providing the necessary infrastructure to handle HTTP requests, render templates, and manage application routing. Its flexibility and minimalistic approach allow for the integration of various extensions and libraries as needed, ensuring the application remains lightweight yet functional. Flask's

comprehensive documentation and supportive community further enhance its suitability by the creation and support of community-driven extensions speeding up development and reducing the need to reinvent the wheel.

However, as development progressed, the need for better I/O performance became increasingly apparent. Early on, it was clear that leveraging Python's native `asyncio` would benefit the project—but a significant portion of the existing codebase, including Flask itself, relied on non-async-compatible libraries. This limitation stemmed from Flask's foundation on WSGI, a synchronous standard developed long before Python's `asyncio` was introduced. WSGI operates strictly in a blocking request-response model, requiring each request to complete fully before processing the next. While traditional workarounds like multi-threading or multi-processing can mitigate this, they introduce complexity and are better suited for CPU-bound tasks than I/O-bound workloads.

To address these constraints, several approaches were considered:

- **Gevent/Eventlet:** These libraries use monkey-patching to emulate asynchronous behavior in synchronous code. However, they are not true `asyncio` and can lead to unpredictable behavior. Given the project's early stage, this option was deemed too risky
- **Flask + Celery:** Offloading long-running tasks to Celery workers helps avoid blocking but introduces operational overhead, requiring additional infrastructure like Redis or RabbitMQ for message brokering, as well as resource overhead, as celery worker pools take up resources even when idle
- **Quart (ASGI Flask):** A Flask-compatible ASGI reimplementation with native `async/await` support. However, Quart lacks Flask's mature ecosystem and still relies partially on monkey-patching, raising concerns about long-term stability
- **FastAPI (Full ASGI migration):** Built on ASGI, FastAPI was designed with `async-first` principles, enabling efficient handling of thousands of concurrent connections. Its native `async/await` support and modern tooling offer a cleaner solution without the need for workarounds, at the expense of having to reimplement some features already implemented in Flask

While Flask remained suitable for early development, emerging requirements—particularly those involving asynchronous communication for more scalable I/O operations—eventually

led to a need to explore architectural shifts, due to Flask's limited async support and WSGI heritage.

4.6.2 Second setup: Flask + Celery

As the limitations of Flask's synchronous WSGI model became more apparent, we explored Celery as a potential solution for handling asynchronous tasks. Celery, a distributed task queue system, allows offloading blocking I/O to separate worker processes. Celery operates by decoupling task execution from the main application flow. When a time-consuming operation is required, Flask dispatches it to a Celery worker via a message broker (typically Redis or RabbitMQ). The worker processes the task asynchronously, while Flask remains free to handle incoming requests. While this approach mitigated some of Flask's blocking I/O issues, it introduced new challenges in complexity and system overhead.

The Celery architecture operates through worker processes that listen to the message broker for tasks. These workers run as independent processes, executing tasks marked with Celery's `@app.task` decorator. The system's concurrent processing capability comes from multiple workers operating in parallel, each handling different tasks from the queue. Tasks are Python functions that are decorated with provided Celery decorators such as `@app.task`, causing them to be registered as Celery tasks within the Celery application. This design is particularly valuable for operations like batch processing or scheduled jobs that would otherwise block Flask's synchronous request handling.

Algorithm 2 Calling a Celery Task and Getting the Result

```

1: from celery import Celery
2:
3: app      =      Celery('tasks',      broker='redis://localhost:6379/0',      back-
      end='redis://localhost:6379/0')
4:
5: @app.task
6: def hello():
7:     return 'hello world'
8:
9: result = hello.delay()
10: print(result.get())

```

To execute a task, a Celery task function must be called using the `delay()` method, which will return a result object. This result object can be used to check the status of the task and to retrieve the result once it is available.

Celery supports horizontal scaling by design, allowing multiple worker pools to run on separate physical or virtual machines. This makes it especially effective for handling growing workloads—for example, processing email newsletters for an expanding user base.

Celery’s advanced features, including task retries, chaining, and prioritization, while powerful, further increased the system’s complexity. We found ourselves managing not just our application logic, but also the reliability of the message broker, persistence of results, and supervision of worker processes. This architectural overhead seemed increasingly disproportionate to our actual needs as the project evolved.

Furthermore, Celery clients and workers introduce a non-negligible overhead in terms of CPU and memory usage, even when idle, as they must maintain persistent connections to the broker and periodically perform health checks or heartbeats. This can be a concern in resource-constrained environments or during development. This overhead became especially evident during early integration tests.

As the project evolved and tests were performed, it became increasingly clear that Celery’s benefits lend themselves better to CPU-bound workloads, as opposed to our I/O-bound ones, and they did not outweigh its resource and architectural costs for the current use case. This realization prompted an exploration of more lightweight asynchronous alternatives, eventually culminating in an investigation into ASGI-compliant frameworks with native async capabilities and simpler concurrency management.

4.6.3 Third setup: Quart, an ASGI-compliant Flask reimplementation

Quart emerged as a promising candidate during our exploration of async solutions, offering a unique combination of Flask syntax with compliance. As a near-drop-in replacement for Flask, Quart theoretically allowed for an easy migration while providing all the benefits of native async/await support. The framework’s design promised seamless execution of asynchronous code alongside familiar Flask patterns, making it particularly attractive for existing Flask projects that would benefit from asynchronous capabilities.

However, practical evaluation revealed significant limitations in Quart’s ecosystem maturity. While the core framework maintained good compatibility with code originally written for Flask and critical extensions we relied on - including authentication, database integration, among others - either lacked Quart equivalents or had poorly maintained implementations. We discovered many Quart-specific packages were either abandoned, documented only through sparse READMEs, or failed to match their Flask counterparts in functionality. For

instance, the Flask-Login equivalent for Quart was one such unmaintained extension that would require some reimplementation. This gap would require us to reimplement substantial portions of our web application instead of using existing, known-good solutions.

While Quart offers the possibility to use monkey patching on Flask extensions, deeper evaluation revealed this compatibility came with significant trade-offs. This approach could make some Flask extensions work in the async environment, we found this to be an unstable foundation for long-term maintenance. Notably, recent Quart releases have started to move away from this approach - the framework now treats these compatibility layers as optional extras rather than core features, signaling a deliberate shift in architectural direction.

Additionally, Quart's smaller community and limited production adoption made it difficult to assess long-term viability, raising concerns about framework maintenance and the availability of future support.

Ultimately, while Quart's technical merits as an ASGI Flask alternative were sound, the ecosystem risks and migration costs outweighed its benefits for our project. The framework's current state appears best suited for teams that can commit to Quart's entire stack, rather than as a migration path for existing Flask applications with extension dependencies. This realization steered us toward more mature ASGI alternatives, despite requiring some reimplementation, that could provide robust async support without sacrificing ecosystem stability.

4.6.4 Final Decision: FastAPI Migration

After thorough evaluation of the previous options, we ultimately selected FastAPI as our framework of choice. The decision was driven by FastAPI's native support, which provides built-in asynchronous capabilities without requiring additional infrastructure components. Unlike our Flask + Celery implementation that demanded separate worker processes and message brokers, FastAPI achieves comparable performance under high concurrency through asyncio-compatible architecture, which enables concurrent non-blocking I/O, while significantly reducing system complexity and operational overhead. The framework's performance characteristics proved particularly advantageous for our I/O-bound operations, matching Celery's throughput in load testing but having a simpler application architecture.

The transition to FastAPI brought multiple technical benefits beyond just asynchronous capabilities. The framework's integrated OpenAPI documentation generation and Pydantic-based data validation significantly improved our development workflow and API reliability. While the migration required adapting our route handlers and dependency management patterns, this investment was offset by FastAPI's excellent documentation and growing ecosystem. FastAPI emerged as the most balanced solution, combining mature production readiness with modern features and approachable syntax. The framework's design addressed our immediate performance requirements but also established a solid foundation for future enhancements.

5. Implementation

This chapter will expand on topics introduced in the last chapter, while going more into implementation specifics and providing more details.

5.1. Project structure

In this section we will go over the structure and technologies used in the implementation of our project.

```
Code
├── app
├── inventory
├── logger
├── nornir_lib
├── gns3_api
└── proxmox_api
```

5.1.1 Technologies

The architecture emphasizes separation of concerns through several key design choices. A modular package structure organizes code into logical components, while the use decorators and dependency injection promotes code reuse. Strict interface boundaries between components maintain clear contracts, and repository patterns abstract data access details. This approach adheres to DRY principles while ensuring maintainability as the project scales.

5.1.1.1 Python

All application components were developed in Python 3.10+, chosen for its mature async/await implementation and robust type system. FastAPI serves as our web framework, having replaced earlier Flask + Celery-based prototypes due to its superior native async support. External API communications are handled through HTTPX for asynchronous HTTP REST interactions, while network device configuration validation is managed via Nornir.

5.1.1.2 SQLite

For our database, we adopted SQLite during development, accessed through SQLAlchemy - an Object-Relational Mapping (ORM) built atop SQLAlchemy that incorporates Pydantic's validation capabilities. This combination provided type-safe queries through Python type hints while maintaining SQLAlchemy's powerful query syntax, along with seamless FastAPI integration for automatic OpenAPI schema generation. The use of SQLAlchemy ensures an easy future transition to production-grade databases like PostgreSQL when needed.

5.1.2 app/

The `app/` directory contains the complete implementation of our web application, organized to promote maintainability and clear separation of concerns. The root level includes several key files that form the application foundation:

- `main.py` - The ASGI entry point to run with any ASGI-compliant server
- `database.py` - Manages database connections and session handling
- `models.py` - Defines all database tables and their relationships using SQLAlchemy
- `decorators.py` - Contains reusable decorators for route handling and logic
- `config.py` - Centralizes application settings loaded from environment variables

You will also find the following folder structure

```

app
├── alembic
├── dependencies
├── repositories
├── routers
├── services
├── templates
├── uploads
└── utils

```

`alembic/` responsible for database creation and migrations, aswell as seeding a small amount of dummy data.

`dependencies/` contains all FastAPI dependency injections, including shared resources like model repositories, along with authentication and authorization utilities. These components are reused across multiple endpoints.

`repositories/` implements the database abstraction layer using `SQLModel/SQLAlchemy`, following the repository pattern. This directory houses all database queries and data access operations, providing a clean interface to access required data.

`routers/` organizes API endpoints by domain (authentication, exercises, vms etc.). Each router file contains related route definitions with minimal logic, delegating complex operations to the services layer.

`services/` forms the core logic layer, processing data between repositories and routers. This directory contains hides the complex workflow of some functions and offers a clean interface.

`templates/` stores Jinja2 templates for front-end interfaces, along with related static assets. The templates follow a consistent layout system.

`uploads/` stores all files uploaded by priviledged users, mainly `.gns3project` files.

`utils/` provides shared utility functions and classes that don't belong to specific domains.

5.1.3 inventory/

The `inventory/` directory contains YAML configuration files that store static device information for each topology. These files serve as a device registry, enabling Nornir to immediately access device specifications without runtime type checking. This approach significantly improves automation performance by eliminating redundant device discovery operations.

Table 5.1: Example inventory file contents

Device	Groups	Hostname	Port	Username	Password
linuxvm-1	linuxvm	192.168.57.143	5005	<username>	*****
pc1	vpcs	192.168.57.143	5007	—	—
r1	cisco_router	192.168.57.143	5000	—	—
sw1	cisco_switch	192.168.57.143	5002	—	—

Key Fields Explanation:

- **Device Name:** Unique identifier within the topology (e.g., `linuxvm-1`, `r1`)
- **Device Class:** Specifies the device type/role (determines connection handlers)
 - `linuxvm`: Linux hosts
 - `cisco_router`: Cisco IOS Router
 - `cisco_switch`: Cisco IOS Switch
 - `vpcs`: Virtual PC simulator
- **Hostname:** Shared IP indicates all devices are virtual instances on the same host
- **Port:** Unique port for each device's management interface
- **Credentials:** Shown here as placeholders

Implementation Notes:

1. The YAML structure enables easy integration with Nornir's inventory plugins
2. Port assignments do not follow any specific order
3. Credential fields use `null` values for unauthenticated devices like VPCS

5.1.4 logger/

The `logger/` directory implements the application's centralized logging system with consistent configuration across all components. This module provides:

- Pre-configured log formatting with timestamps, log levels, and module names
- Simultaneous output to both file (`app.log`) and console
- Easy integration via `get_logger()` factory function

The logger module enforces standardized logging practices across the entire application, featuring a consistent log format that includes timestamps, severity levels, and module identifiers for all log entries. Configuration is centralized in `logging_config.py`, which automatically creates and manages the `app.log` file in the project root directory while

simultaneously outputting to the console. By default, the system logs messages at INFO level and above, with built-in flexibility to adjust verbosity for debugging purposes through simple configuration changes. This approach ensures uniform logging behavior while maintaining adaptability for different runtime environments.

The implementation follows Python best practices while allowing for future extensions such as log rotation or remote logging services. All application modules should obtain their logger instance through the provided `get_logger()` function to maintain consistent logging behavior.

5.1.5 `nornir_lib/`

The `nornir_lib/` directory implements the evaluation system that interfaces with various virtual devices. This component executes commands across network topologies and analyzes the responses to determine operation success.

Configuration requires three key files in the `app/` directory:

- `config.yaml` - Specifies paths to host/group files and Nornir runner settings
- `host_file` - Contains device credentials (IP, username, password in plaintext)
- `group_file` - Defines device group parameters (must maintain `fast_cli: false`)

Developers can implement new test modules by extending the base `CommandModule` class.

This requires implementing device-specific command methods (`_command_router()`, `_command_switch()`, etc.) and corresponding response interpreters (`interpret_cisco_response()`, etc.) with the benefit of not having to worry about anything about nornir and its inventory system. The system currently includes Ping and Traceroute implementations, with the modular design. Developers can use the bundled Ping module that demonstrates this pattern, taking parameters and evaluating success based on configurable packet loss tolerance.

The architecture emphasizes:

- Consistent device communication through standardized interfaces
- Flexible test creation via module inheritance
- Centralized response interpretation logic

In the past iteration of this project, it was noted that communication with devices could be less than reliable, with no apparent reason. After some testing and research it was found that disabling the `fast_cli` increased reliability and we have not experienced the communication failure since disabling this feature.

5.1.6 proxmox_api/

The `proxmox_api` library provides direct, lightweight wrappers around the Proxmox VE REST API, offering simplified interfaces for common virtualization management tasks while maintaining close correspondence with the underlying API endpoints.

Design Philosophy

- **Transparent Wrapping:** Each method maps clearly to a specific Proxmox VE API endpoint
- **Minimal Abstraction:** Preserves the API's native behavior with light conveniences
- **Consistent Error Handling:** Uniform approach across all operations
- **Asynchronous Communication:** Use async methods where possible to maximize performance

5.1.6.1 Error Handling

The library implements a consistent error handling approach through the `@handle_network_errors` decorator, which manages network-related exceptions while preserving application-level errors. This decorator specifically intercepts connectivity issues (host unreachable, time-outs) and HTTP 404 responses, converting them to `False` returns while maintaining detailed error logging. All other HTTP errors and programming exceptions propagate unchanged, ensuring callers receive complete error context for non-network failures.

Table 5.2: Error Handling Behavior

Case	Behavior
Network Connectivity	Returns <code>False</code> with error logging.
HTTP 404 (Not Found)	Logs and re-raises with request details.
Other HTTP Errors	Propagates with original status code.
Application Exceptions	Unmodified propagation.

The implementation preserves function signatures through Python's `@wraps` decorator and maintains type safety via generic type variables. Designed specifically for async operations, it provides transparent error handling that distinguishes between temporary network issues and substantive application errors while ensuring comprehensive diagnostic logging.

Implementation Patterns The methods follow three primary patterns:

Table 5.3: Method Implementation Patterns

Pattern	Characteristics
Simple Wrapper	Single API call + status check (e.g., <code>start</code> , <code>stop</code>)
Chained Operation	Multiple API calls (e.g., <code>create</code>)
Special Case Handler	Custom status code processing (e.g., <code>check_free_id</code>)

proxmox_vm_actions Module Key VM operations:

- `start(proxmox_host, session, vm_id)`
 - Issues `POST /nodes/<node>/qemu/<vmid>/status/start`
 - Verifies successful status change
- `create(proxmox_host, session, template_id, clone_id, hostname)`
 - Chains clone operation + protection removal
 - Handles storage and naming configuration

5.1.7 gns3_api/

The GNS3 API wrapper provides essential operations for managing GNS3 projects through its REST API, following similar design patterns to the `/proxmox_api` wrapper but tailored for GNS3-specific workflows. The library handles project operations including verification, import/export, and node information collection.

Implementation Characteristics The wrapper shares several architectural features with the `proxmox_api` implementation:

- Uses the same `@handle_network_errors` decorator pattern
- Follows consistent `async/await` patterns
- Maintains similar logging practices

However, it differs in several GNS3-specific aspects:

- Project-centric operations rather than VM management
- File handling for project import/export
- UUID-based project identification

Error Handling The library maintains consistent error behavior:

- Returns `False` for network failures
- Returns `None` for missing resources
- Propagates all other exceptions

Example Workflow A typical usage sequence would be:

1. Verify project exists (`acheck_project`)
2. Get node information (`aget_project_nodes`)
3. Start project nodes (`astart_project`)

The implementation demonstrates careful handling of both network operations and local file I/O, particularly in the import/export methods where it manages binary data transfer and local filesystem interactions. The UUID-based project identification ensures unique project references during import operations.

5.2. Web Application Components

This web application is structured to offer the following capabilities:

1. Login with institutional or local credentials

2. Selection from available exercises
3. Automated environment preparation
4. Practical work in GNS3 web interface
5. Validation feedback

To accomplish that, it comprises three core modules that work in concert to manage net-working exercises while abstracting the underlying virtualization infrastructure:

5.2.1 Authentication Module

The authentication module establishes user identity and access control through JWT to-kens. It verifies ownership of virtual resources during every operation, preventing cross-user interference like unauthorized VM control. To help accomplish these features it inte-grates with the database to check for priviledged accounts and VM ownership.

5.2.2 Exercise Management Module

This module handles the complete exercise workflow from creation to validation. Instruc-tors can upload network topologies, by providing gns3project files and validation criteria using the previously mentioned validation modules, namely Ping and Traceroute, during exercise creation, while students receive filtered exercise lists based on which they are enrolled in. The validation subsystem uses the developed modules to validate instruc-tor defined criteria, providing automated feedback for students. All provisioning occurs automatically when students are enrolled in exercises.

5.2.3 VM Control Module

Finally this modules provides the ability to interact with VMs by exposing endpoints for, among other things, powering on/off and request exercise validation. This allows students to avoid interacting directly with the underlying infrastructure and focus on doing their exercises.

5.3. Asynchronous Processing with FastAPI

Asyncio is Python library for writing concurrent code. It provides a foundation for asyn-chronous programming by enabling the creation and management of event loops, corou-tines, and asynchronous tasks.

An *event loop* is a central component of asynchronous programming—it continuously runs in the background, managing the execution of asynchronous tasks. When a task reaches a point where it would normally block (e.g., waiting for a network response), it yields control back to the event loop, which can then continue running other ready tasks. This model of cooperative multitasking contrasts with traditional multithreading or multiprocessing, as it operates in a single thread and does not require locking or context switching between OS threads.

A *coroutine* is a special kind of function defined with `async def`. When called, it does not run immediately, but instead returns a coroutine object. This object can be scheduled by the event loop, and when awaited, it runs until it hits a pause point (e.g., another `await`)—at which point it yields control back to the event loop, allowing other coroutines to execute.

In FastAPI, declaring an endpoint as `async def` enables non-blocking behavior for I/O operations when using async-compatible libraries. This allows the server to handle other requests while waiting for operations like external API calls. If that logic includes `asyncio`-compatible I/O operations—such as using a library for asynchronous HTTP calls then the request can proceed in a truly asynchronous manner. This allows the web server to observe massive speedups when compared to blocking I/O when multiple HTTP calls must be made to external services.

Additionally, `asyncio` supports the orchestration of multiple tasks using constructs such as `asyncio.gather()`, which allows multiple coroutines to be executed concurrently and awaited collectively. This has been especially useful in scenarios within the project where multiple devices or services must be queried or configured simultaneously, such as when multiples students are enrolled in an exercise, which requires the creation of multiple VMs.

5.3.1 Differences Between Asyncio And Celery

In contrast, Celery operates at a higher level of abstraction, focusing on distributed task execution rather than fine-grained concurrency. Instead of relying on an event loop, Celery uses a pool of worker processes—often distributed across multiple machines—that consume tasks from a message broker such as RabbitMQ or Redis. Tasks in Celery are standard Python functions decorated with `@app.task`, which serializes their execution requests into messages sent to the broker. Workers then fetch these messages and execute the tasks in separate processes, enabling true parallelism across CPU cores or even different servers. Celery’s architecture makes it particularly well-suited for workloads that

require heavy computation, long-running operations, or distributed execution across multiple machines. Unlike `asyncio`, which excels at managing many lightweight I/O-bound tasks within a single thread.

While both `asyncio` and Celery outperform traditional sequential blocking I/O code, `asyncio` proved better aligned with our project's requirements. Sequential code suffers from inherent inefficiencies: each I/O operation forces the program to idle while waiting for a response, wasting CPU cycles that could be used for other tasks. `asyncio` eliminates this waste by allowing multiple of I/O operations to proceed concurrently within a single thread, dramatically improving throughput for I/O-bound workloads. Celery, while also avoiding blocking behavior, introduces overhead from inter-process communication and task serialization, making it less optimal for high-frequency, low-latency operations. In our use case, where the system primarily handles short-lived HTTP requests, `asyncio`'s lightweight coroutines delivered the same or even superior performance with simpler structure and code. Celery remains invaluable for projects requiring background jobs, but for real-time, I/O-heavy scenarios, `asyncio` provided the speed we wanted with efficient resource usage and also being more maintainable.

5.4. GNS3 Customization and Configuration

This section outlines the configuration process for a GNS3 host virtual machine. This setup must be performed during the initial deployment of the system, as it involves creating a base template VM that includes a properly configured GNS3 instance. All future clones used by the system will be derived from this base template.

The first step involves installing the `gns3-server` along with all its required dependencies. This provides the core backend functionality. The installation can be done using a provided remote installation script that handles the setup of Python packages, IOU support, and necessary architecture extensions such as the i386 repository. This script can be found on the official GNS3 website.

Once installed, it is essential to configure the `gns3-server` to run as a system daemon. Running the server as a background service, ensuring it is automatically started at boot time and restarted in case of crashes, remaining continuously available without requiring manual intervention. This is especially important to ensure no manual interaction is needed with the host VM.

In addition to installation and service configuration, modifications must be made to the `gns3-server` source code to enable an auxiliary telnet console port while keeping SPICE support. SPICE is a remote display protocol that allows users to interact with VMs through a graphical interface. These source code changes enable the server to launch QEMU instances with the appropriate SPICE options, facilitating enhanced remote access and control over the virtualized devices. A more detailed explanation and rationale for these changes can be found in the first iteration of this project[1].

The host operating system for the GNS3 host VMs during development was **Ubuntu Server** (non-minimized installation). This ensures that all necessary system tools and dependencies are available. Other operating systems such as different Linux distributions or Windows were not tested.

Note: The `gns3-web` UI for GNS3 is currently in beta and may present issues when adding templates for devices. At this point in time it is recommended to use the GUI client for this task.

Note: For clients to interact with SPICE enabled devices, such as virtualized linux hosts the user must install `gns3-webclient-pack` on their machine.

IOU Support: If the project includes IOU nodes, a valid IOU license is required. This file must be placed in `~/.iourc` and formatted according to GNS3's expectations.

5.5. Proxmox API Usage

Proxmox VE provides a REST API that exposes all functionality available. This includes operations such as creating, cloning, starting, stopping, and deleting virtual machines, as well as querying their current status. By interfacing with this API, the system gains the ability to manage VMs in an automated, repeatable, and scalable manner, which is essential for deploying work environments on demand.

In our project, the Proxmox VE API is accessed via HTTP using the HTTPX library in Python. Authentication is handled using Ticket Cookies. A ticket is a signed random text value with the user and creation time included. Additionally, any write (POST/PUT/DELETE) request must include a CSRF prevention. To obtain a valid ticket and CSRF token a POST request must be made to the appropriate endpoint with valid plaintext credentials in the

body of the message. This ticket is valid for a set amount of time and itself can be sent to the correct endpoint to acquire a fresh ticket.

To avoid repeated logins, aswell as avoiding storage of plaintext user credentials, and reduce overhead, tickets are stored in memory and reused for their duration, after which a new token is acquired.

Additionally, when users log into the web application, their provided credentials are sent to the Proxmox VE API authentication endpoint to obtain a ticket specific to that user. This approach offers multiple advantages. First, regular user accounts, such as those belonging to students, can be configured with limited permissions in Proxmox VE. This serves as a secondary layer of protection in case of an exploit in the authentication logic that allows an unauthorized request (e.g., attempting to delete a VM). Second, it improves accountability, as every action taken by a user can be traced to their authenticated account, resulting in more detailed and accurate audit logs.

For this approach to work, valid user accounts must exist within the Proxmox VE system. These accounts can be created in several ways, as Proxmox VE supports a variety of authentication backends. For example, integration with LDAP is supported out of the box. This makes user management particularly convenient in educational institutions, where centralized identity systems are commonly used. Alternatively, local accounts managed directly within Proxmox VE can also be used, offering a simpler setup for smaller-scale or isolated environments. The use of local accounts or directory services for Proxmox VE should go in tandem with the use in the web application.

6. Testing & Evaluation

In this chapter we will go over how we did performance and stress tests as well as discuss the results seen.

6.1. Performance Evaluation

To evaluate the impact of migrating from Flask (+ Celery) to FastAPI, we conducted a series of measurements focused on two main performance metrics:

- **Time to complete a given task.**
- **CPU resources consumption.**

Each test consisted of three sequential tasks:

1. **1st task: Template VM creation:** Create a linked clone from a pre-configured template. Once the cloning process is complete, the VM is powered on and polled periodically until it obtains a valid IP address. A project is then imported into its GNS3 instance. After the import completes, the VM is powered off and converted into a reusable template VM.
2. **2nd task: VM cloning:** Generate a specified number of linked clones from an existing exercise template, each intended for individual user assignments.
3. **3rd task: VM deletion:** Remove all cloned VMs associated with an exercise, as well as the corresponding template VM.

Each batch test was conducted using different quantities of VM clones: 1, 10, 20, 100, and 200.

All three implementations were deployed in the same network environment and executed on identical hardware to ensure a fair comparison. The web application in each case was hosted within an LXC container running on the same Proxmox VE machine. The container was allocated 8 virtual CPU cores and 1GB of RAM.

Time measurements were recorded from the start of a task's execution, defined as the moment the first relevant API request was issued, until the moment of successful completion, when the final expected API response was received.

Each test scenario was repeated three times per implementation. The reported values represent the average across these repetitions.

For the FastAPI implementation, a concurrency limit of 5 was imposed for interactions with the Proxmox VE API. Other categories of tasks, such as user requests, GNS3 project handling, and Nornir automation, were allowed to proceed without any concurrency restrictions.

6.1.0.1 Task 1: Template VM Creation

This task is sequential in nature, as it interacts with only a single VM, and none of the steps can be performed in parallel. As such, no significant performance uplift is expected.

Table 6.1: Comparison of Execution Time (in seconds)

Task	Flask	Flask + Celery	FastAPI
Template VM creation	35.745697	34.639572	32.202087

As expected, the average execution times are very similar across all three implementations, with the FastAPI approach showing a slight performance advantage.

However, averages alone do not tell the full story. Figure ?? illustrates the execution time variability. *Note: the y-axis does not start at 0.*

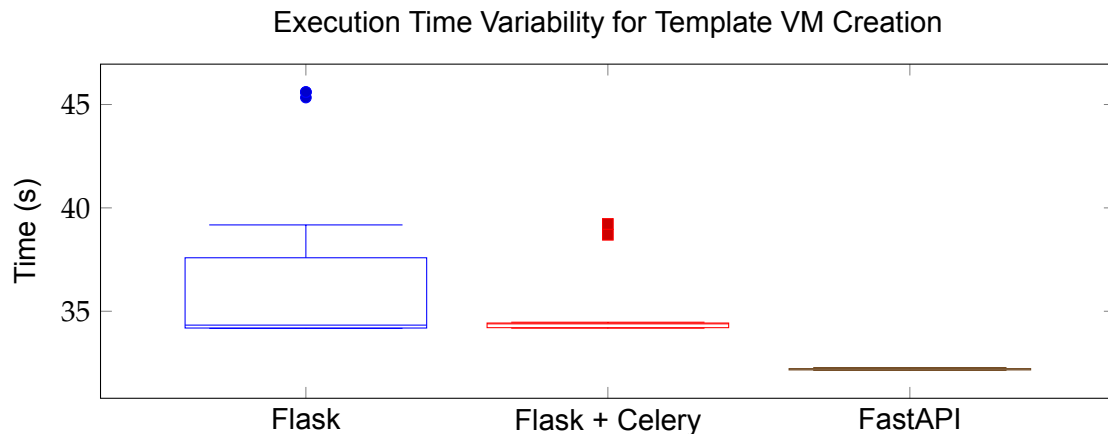


Figure 6.1: Boxplot comparison of execution time variability for Template VM creation.

From the boxplot, it is evident that the Flask implementation exhibits the highest variability in execution times, indicating inconsistency. The Flask + Celery variant improves consistency but still contains several statistical outliers. In contrast, the FastAPI implementation demonstrates significantly tighter grouping and no apparent outliers, making it the most consistent of the three.

6.1.0.2 Task 2: VM Cloning

Unlike the previous task, the cloning of multiple VMs can be executed concurrently, as each cloning operation consists of only two main steps and is entirely independent from others.

First, a random VM ID is generated, followed by a check to ensure that the ID is not already in use. This typically requires just a single API call and succeeds immediately in nearly all cases, given the extremely low probability of collision within our ID range (100 to 999999999).

Once a free ID is found, the clone is created using that ID. Since these steps are isolated per VM, a large number of cloning operations can be run in parallel with minimal contention or coordination overhead.




Figure 6.2: Execution time to clone increasing number of VMs across different implementations.

6.1.0.3 Task 3: VM Deleting

6.1.1 I/O Problems during Batch VM Operations

During stress tests of mass cloning and deletion of VMs (200 at a time, dispatched to a single Proxmox VE node), we observed rising task times, specifically for the mass cloning of new VMs, as can be seen in the following graph

It was also noted that there was an error during the 4th batch of tasks, specifically during the deleting task.




Figure 6.3: Grouped bar chart showing VM operation times. Cloning time is labeled to highlight the rising trend. Results from Flask running purely sequential code

After some investigation it was found that there were intermittent failures to remove VM disks. These failures were *not* detectable via the Proxmox VE HTTP API responses and only appeared in the Proxmox VE server logs. As orphaned disks accumulated, overall performance degraded significantly.

Observed Task History Outputs:

--- 1st type of output (VM disks removed successfully) ---

trying to acquire lock...

OK

Logical volume "vm-348786940-disk-0" successfully removed.

TASK OK

--- 2nd type of output (intermittent lock-timeout failures) ---

trying to acquire lock...

Could not remove disk 'local-lvm:vm-120993831-disk-0', check manually:

can't lock file '/var/lock/pve-manager/pve-storage-local-lvm' - got timeout

trying to acquire lock...

OK

Logical volume "vm-120993831-disk-0" successfully removed.

TASK OK

--- 3rd type of output (persistent failures) ---

trying to acquire lock...

Could not remove disk 'local-lvm:vm-363495383-disk-0', check manually:

can't lock file '/var/lock/pve-manager/pve-storage-local-lvm' - got timeout

trying to acquire lock...

can't lock file '/var/lock/pve-manager/pve-storage-local-lvm' - got timeout

TASK OK

--- 4th and final type of output (storage config update errors) ---

trying to acquire lock...

Could not remove disk 'local-lvm:vm-5469324-disk-0', check manually:

can't lock file '/var/lock/pve-manager/pve-storage-local-lvm' - got timeout

trying to acquire lock...

can't lock file '/var/lock/pve-manager/pve-storage-local-lvm' - got timeout

trying to acquire cfs lock 'file-user_cfg' ...

TASK OK

This suggests Proxmox VE's locking mechanism can't keep pace with big amounts of quick sucession delete requests. The lock is used to ensure that two tasks dont modify the LVM's metadata simulatenously. Since there should be some background tasks that are piling up, the system cant keep pace and eventually starts using retries to keep up but even this is insuficient and starts failing more and more towards the end as it becomes fully congested. At the end the system starts becoming overwhelmed and also begins having trouble updating its internal storage config file.

This, combined with other factors were the main reason that led us to implement a hard limit on the amount of concurrent requests that can be made from the web application to Proxmox VE API. Still, while this significantly reduces the chances of this problem reoccurring, it does not fully remedy the problem and additional future work should look into solving this matter completely.

After performing a cleanup of the orphaned disks we can see that performance was improved from even the 1st baseline test.




Figure 6.4: Bar chart showing VM cloning times after cleanup, at baseline, and in the worst case scenario.

7. Conclusion & Future Work

References

- [1] P. M. C. Santos, “Environment for practical evaluations in network administration,” Master’s thesis, Faculdade de Ciências da Universidade do Porto, Porto, Portugal, 2024, accessed: March 21, 2025. [Online]. Available: <https://hdl.handle.net/10216/164899> [Cited on pages 2, 15, and 50.]
- [2] GNS3 Documentation, “Architecture,” n.d., accessed: March 15, 2025. [Online]. Available: <https://docs.gns3.com/docs/using-gns3/design/architecture/> [Cited on page 5.]
- [3] Proxmox Server Solutions GmbH, *Proxmox VE Administration Guide*, 2025, accessed: March 21, 2025. [Online]. Available: <https://pve.proxmox.com/pve-docs/pve-admin-guide.html> [Cited on page 6.]
- [4] P. J. Eby, *PEP 333 – Python Web Server Gateway Interface v1.0*, Online, 2003, accessed: March 21, 2025. [Online]. Available: <https://peps.python.org/pep-0333/> [Cited on page 10.]
- [5] P. Projects, “Flask documentation,” Online, 2025, accessed: March 21, 2025. [Online]. Available: <https://flask.palletsprojects.com/en/stable/> [Cited on page 11.]
- [6] D. Barroso, “Nornir: The python automation framework,” Online, 2025, accessed: March 21, 2025. [Online]. Available: <https://nornir.readthedocs.io/en/latest/> [Cited on page 14.]
- [7] J. Leal and F. Silva, “Mooshak: a web-based multi-site programming contest system,” *SOFTWARE-PRACTICE & EXPERIENCE*, vol. 33, no. 6, pp. 567–581, MAY 2003. [Cited on page 17.]

Appendix Title Here

Write your Appendix content here.